IEEE **HPEC** 2023

# Spla: Generalized Sparse Linear Algebra Library With Vendor-Agnostic GPUs Acceleration

Egor Orachev<sup>1</sup> Semyon Grigorev<sup>1</sup>

egor.orachev@gmail.com s.v.grigoriev@spbu.ru, <sup>1</sup>St. Petersburg University, Russia



### Problem Statement

Scalable high-performance graph analysis is a challenge. GraphBLAS standard attempts to solve this challenge using sparse linear algebra operations. The full GPU-based implementation of this standard is still missing. Existing works are focused on Nvidia Cuda platform only, what limits their potability.

it possible to implement portablelibraryGPU-based  $with \quad generalized$ sparse linear algebra operations for graph analysis?

### Results

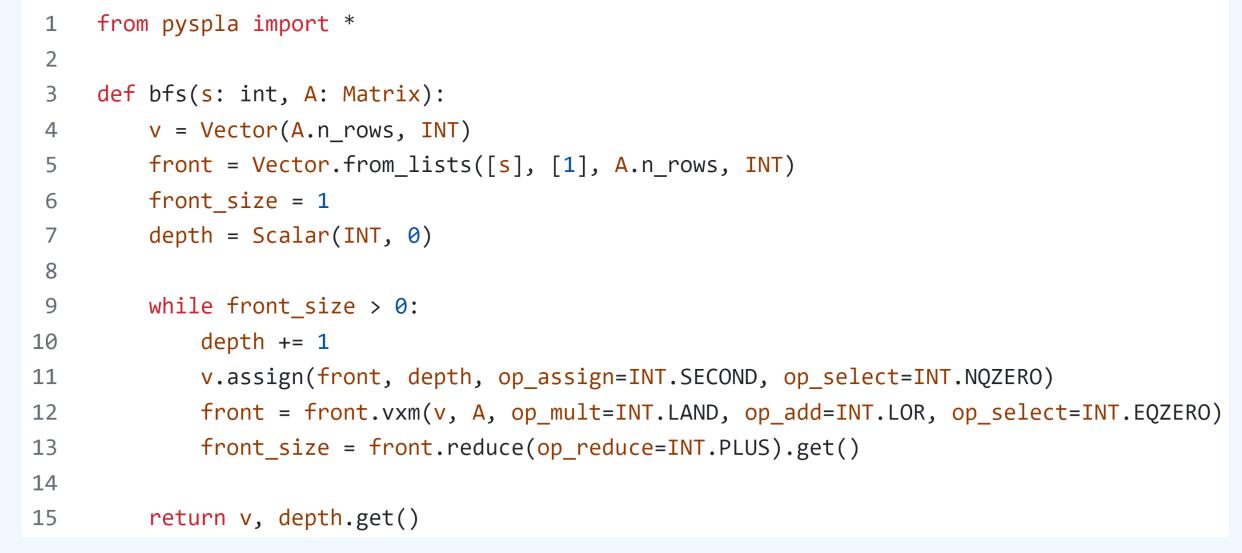
Spla, the generalized sparse linear algebra library with vendor-agnostic GPUs accelerated computations.

- Proven **portability** of the OpenCL-based solution.
- Demonstrated **competitive** performance.
- Had acceptable **scalability** on devices of different GPUs vendors.
- Published **package** for python users.

### Future Research

- Adaptive workload balance for better GPU's waves utilization.
- Performance tuning of auxiliary utilities, such as sorting, reduction, etc.
- New operations variations, such as SpGEMM, SpMM, etc.
- Graph data from RAM to VRAM streaming.
- Schalability to multiple GPUs.

### Motivating Example



Push-only BFS implementation with out of the box GPUs acceleration using pyspla package API available at PyPI.

### Implementation

Spla implemented using the C++17 language and vendoragnostic OpenCL 1.2 for GPU specific computations. General optimizations.

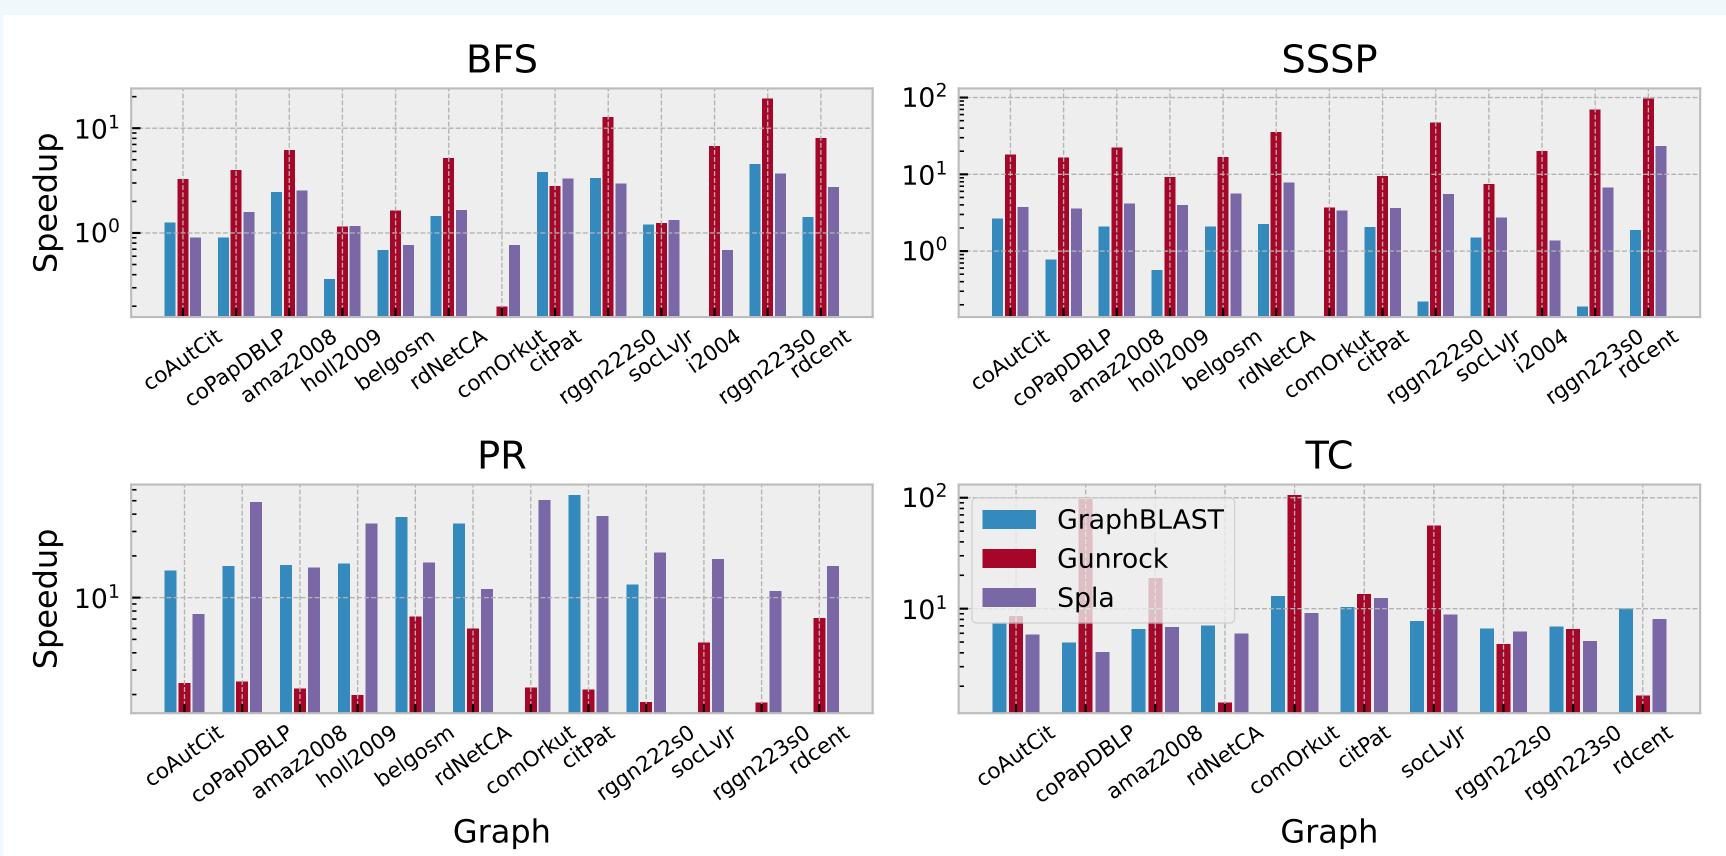
- K-way merge-based vxm operation [1].
- Push-pull and early-out [2].
- Masking of reached vertices [3].
- Sparse-dense storage switch [3].

### OpenCL Specifics

Auxiliary functionality and optimizations to reduce mostly OpenCL CPU-side driver overhead.

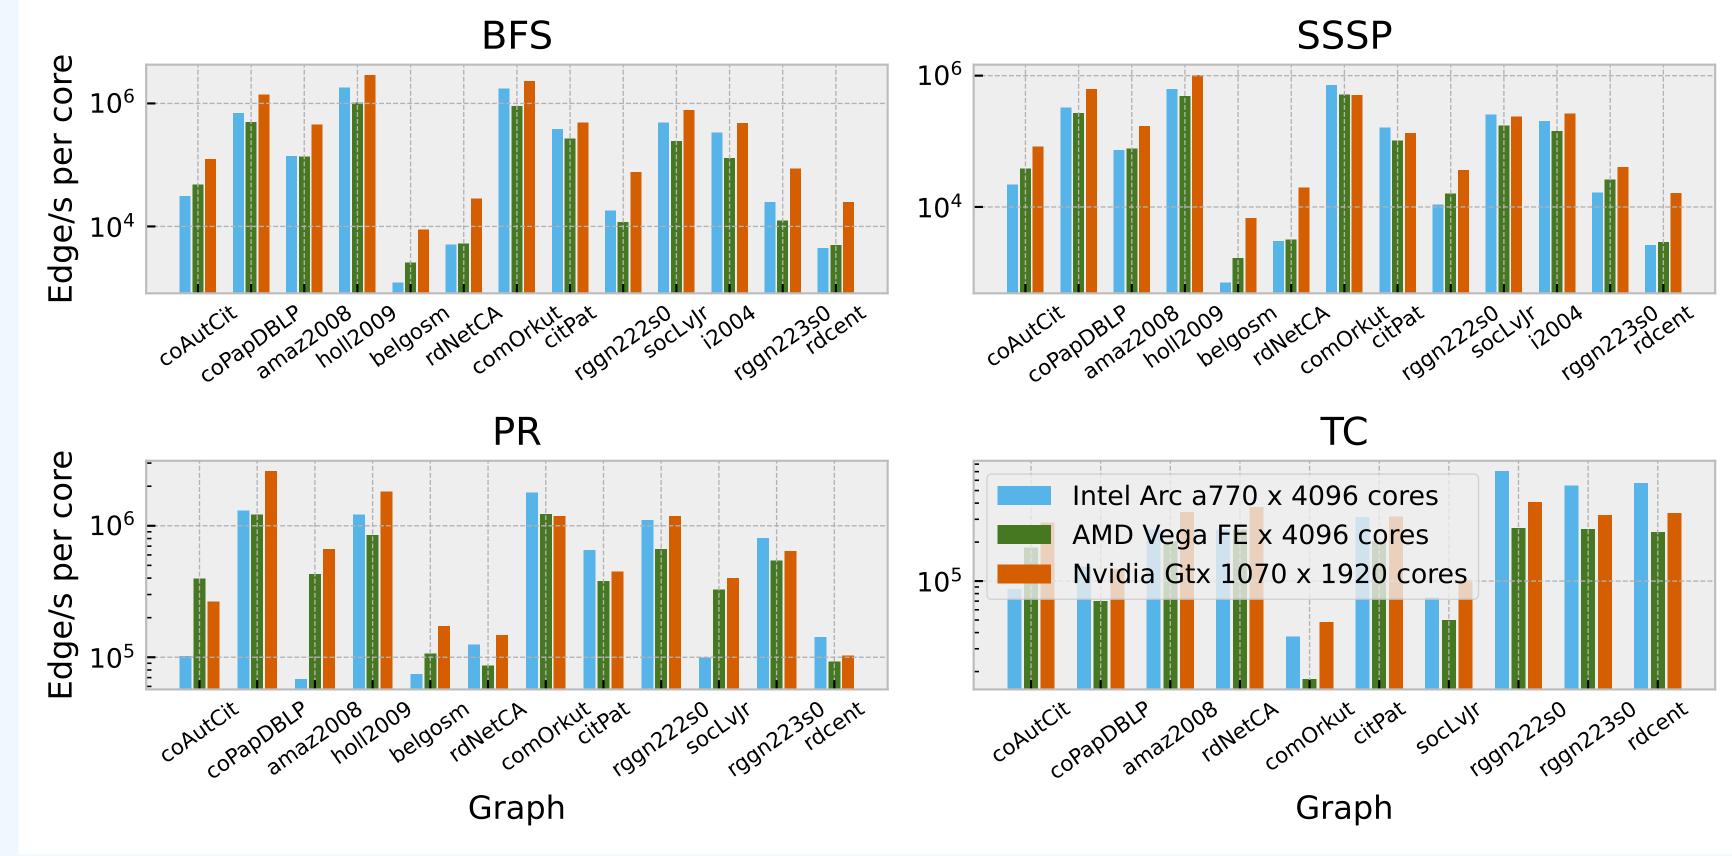
- Run-time compiled **kernels cache** with a robin hood hash map for fast look-ups.
- In-place unique kernel key string constriction.
- Kernel code compilation with **user-defined** functions using simple text preprocessing.
- Custom linear allocator based on sub-buffer mechanism for temporary small ( $\leq 1MiB$ ) allocations.
- Custom auxiliary operations, such as sort, reduce, scan.

### Evaluation: Performance on a Nvidia GPU



**Tools:** Gunrock [4], GraphBLAST [3], Spla and LaGraph [5] as a baseline. Configuration: Ubuntu 20.04, 3.40Hz Intel Core i7-6700 4-core CPU, DDR4 64Gb RAM ♦ Nvidia GeForce GTX 1070 GPU with 8Gb VRAM.

## Evaluation: Scalability on different GPUs



Configuration: Ubuntu 20.04, 3.40Hz Intel Core i7-6700 4-core CPU, DDR4 64Gb RAM ♦ Nvidia GeForce GTX 1070 GPU 8Gb VRAM ♦ Intel Arc A770 flux GPU 8GB VRAM ♦ AMD Radeon Vega Frontier Edition GPU 16GB VRAM.

#### Contact Us

### Our team:

- Semyon Grigorev: s.v.grigoriev@spbu.ru
- Egor Orachev: egor.orachev@gmail.com

### Acknowledgments

We would like to thank Huawei Technologies Co., Ltd for supporting this research. We are also grateful to our team and reviewers.

#### References

- [1] Carl Yang, Yangzihao Wang, and John D. Owens. Fast sparse matrix and sparse vector multiplication algorithm on the gpu. In 2015 IEEE International Parallel and Distributed Processing Symposium Workshop, 2015.
- [2] Carl Yang, Aydin Buluc, and John D. Owens. Implementing push-pull efficiently in graphblas, 2018.
- [3] Carl Yang, Aydin Buluc, and John D. Owens. Graphblast: A high-performance linear algebra-based graph framework on the gpu, 2019.
- [4] Yuechao Pan, Yangzihao Wang, Yuduo Wu, Carl Yang, and John D. Owens. Multi-gpu graph analytics. In 2017 IEEE International Parallel and Distributed Processing Symposium (IPDPS), 2017.
- [5] Gábor Szárnyas, David A. Bader, Timothy A. Davis, James Kitchen, Timothy G. Mattson, Scott McMillan, and Erik Welch. Lagraph: Linear algebra, network analysis libraries, and the study of graph algorithms, 2021.